



STEPPE

Governable Spaces

Emergent strategy

REDMOND • CENTRAL OREGON • EST. 2026

v1 • 2026-07-09

A member-governed Oregon public benefit nonprofit.

Steppe against two doctrines: Governable Spaces & Emergent Strategy

v1 • July 2026 • living document. An honest scorecard of how Steppe’s shipped system and canon hold up against Nathan Schneider’s *Governable Spaces* (2024) and adrienne maree brown’s *Emergent Strategy* (2017), with the policy changes that follow. Adopted changes are marked **ADOPTED**; anything touching member governance is a **PROPOSAL** for the cohort/board — consistent with the standing rule that governance policy is ratified by members, not decided in a commit.

Evidence cites the actual system (schema, app, canon docs), not intentions.

Part I --- Governable Spaces (Schneider)

Schneider’s core diagnosis: online spaces default to **implicit feudalism** — the admin-owns-the-castle pattern baked into tooling (arbitrary founder/mod power, no appeal, no exit with your belongings, rules changeable by fiat). His remedy: **democratic mediums** — spaces whose governance is explicit, modular, learnable, and changeable by the governed.

Where Steppe already embodies it

Doctrine element	Steppe’s implementation	Evidence
Community ownership (no landlord)	Member-governed Oregon nonprofit; can’t be sold; mission entrenched in articles	README; BP v12; NOTICE
Rules-about-rules (metagovernance)	Two-layer constitution: non-votable legal floor vs member-amendable Schedule of Defaults; immutable changes need 75% + 30-day deliberation	Terms §2; CLAUDE.md governance numbers
No fiat rule-change	Governance numbers are provisional config awaiting cohort ratification — the founder deliberately cannot hardcode them as law	CLAUDE.md “Open items”
Due process, not admin whim	Human-in-the-loop on every consequential act; appeals resolved by a different moderator (separation of duties enforced in-DB); every action on a public transparency log; append-only record binding every role incl. service_role	schema.sql:664-678; migration 0012; /protected/transparency
Secret ballot / political privacy	votes readable only by the voter — no moderator clause, no admin path; results only post-close, withheld below a 5-ballot floor	schema.sql:1112-13, 954-974
Exit with your belongings	Full data export (/protected/account/export) + self-service deletion with trust-field scrub	account/export route; delete_my_account()

Doctrine element	Steppe's implementation	Evidence
Anti-capture plumbing	Three-moderator minimum on staggered one-year terms “to preserve plural judgment and prevent capture”; moderators janitorial-not-gatekeeping on taxonomy	Charter 8.9; Groups spec v2 §99
Democratic apprenticeship	Governance is a core app tab, not a settings page; the launch plan's center of gravity is a real vote in week one ; plain language everywhere	GTM Phase 1; invariant P22
Governable stack (infrastructure)	Codeberg-canonical forge, portability audit, year-two self-host milestone — the space is being made governable down to the metal	docs/decisions/codeberg-migration.md; portability-audit.md

This is an unusually strong scorecard; most of Schneider's book is a critique of defaults Steppe was explicitly designed against.

Where feudal residue remains (the honest part)

1. **Moderator selection & recall are unspecified.** Charter 8.9 gives term structure but never says **who appoints, who renews, and how members remove** a moderator. Silence here defaults to appointment-from-above — the exact feudal residue Schneider names. → **PROPOSAL P1** below. (*Highest-priority change in this assessment.*)
2. **The founder-protection period is a disclosed, sunsetted autocracy.** As a transitional device with a constitutional expiry (5–7 yr) it is defensible — Schneider allows for stewardship phases — but its *visibility to members inside the product* is zero today. → **PROPOSAL P6.**
3. **Tenure weighting risks a seniority aristocracy.** 1×–3× voting weight by tenure is live in code (flat **in effect** at launch — a fresh cohort is uniformly 1.0×) and diverges automatically as tenure accrues. The app's own copy promises “members can later vote to weight by tenure” — so the ratification vote is load-bearing, and silence would activate weighting by default. → **PROPOSAL P2.**
4. **Moderation is remedial, not restorative.** The pipeline is report → remove → appeal — procedurally fair but shaped like a courtroom. Schneider (with Ruha Benjamin and the restorative-justice literature) argues democratic spaces need a **repair path**, not only a verdict path. → **PROPOSAL P3.** Deliberately *not* added to the Terms now: promise ≤ practice; build it first.
5. **Rules aren't yet visible as governable objects.** The Schedule of Defaults lives in documents, not in the interface; a member can't see “this rule, amendable, last changed by vote X.” Modular politics wants the rules surface itself legible. → **PROPOSAL P5** (post-launch build).

Part II --- Emergent Strategy (brown)

brown's core principles: **small is good, small is all** (the large is a fractal of the small); **move at the speed of trust** (critical connections over critical mass); **change is constant — adapt intentionally; trust the people; what you pay attention to grows.**

Where Steppe already embodies it

Principle	Steppe's implementation	Evidence
Small is all / fractal	The 35-neighborhood captain structure mirrors whole-org governance; the 12-community hard cap refuses scale beyond governability; “ship the smallest thing that is unmistakably Steppe”	GTM Phase 2; BP v12; CLAUDE.md
Speed of trust / critical connections	Growth only by vouching; cohort “chosen, not solicited”; no paid acquisition ever; partners bring ten people they’d vouch for	GTM non-negotiables + Phase 0
What you pay attention to grows	Health signals are participation and trust metrics, never engagement; chronological + proximity ordering is a constitutional invariant (no attention optimization to grow)	GTM health signals; invariant 7
Intentional adaptation	Living documents with changelogs and revisit clauses; provisional numbers awaiting ratification; DECISIONS.md logs <i>why</i> with a revisit path	DECISIONS.md; BP v12 header
Trust the people (post-gate)	Verify once at the gate, then delete the proof and govern by pseudonymous standing — surveillance-free trust	T&S decision record; verify-then-forget
Interdependence over dependence	Founder-protection sunsets; “an institution that survives its founder”; groups run by their own maintainers	GTM Phase 4; groups RPCs
Guarding emergence against structurelessness	The captains model explicitly imports Jo Freeman’s caution — light, explicit structure so informal power can’t concentrate	GTM Phase 2

The GTM plan is, in practice, an emergent-strategy document that never cites her.

Where it falls short

- Adaptation has documents but no ritual.** Living docs adapt when someone edits them; brown’s discipline is a *rhythm* of collective reflection. The dry-run runbook tests function, not felt experience; nothing schedules the question “what is this like to live in, and what should change?” → **ADOPTED A1** below (an operating practice, not a governance rule — safe to adopt).
- The hardship waiver’s mechanics are unspecified.** “Trust the people” (and they become trustworthy) argues for **self-attestation** — no means-testing theater for a \$4 waiver. Unspecified today; specifying it trust-first is a dues-policy call. → **PROPOSAL P4**.

Part III --- Changes

Adopted now (doc-level, within remit) --- ADOPTED

- **A1 • A reflection cadence for the beta.** During the closed beta: a monthly three-question member pulse (what's working / what's hard / what should change), summarized publicly on the transparency surface; one quarterly retrospective whose outputs are framed **as draft proposals** into the Govern tab — so reflection feeds the governance loop instead of a suggestion box. Cheap, reversible, no code required to start (a pinned listing + a form suffices). Recorded in DECISIONS.md.
- **A2 • Canon acknowledgment.** Schneider and brown join the named intellectual canon (README), so the next reader knows these lenses are load-bearing, not decorative.
- **A3 • This assessment** as a living scorecard, revisited at each stage gate.

Proposals for cohort/board ratification --- PROPOSAL (not decided here)

- **P1 • Moderator selection & recall (fills the Charter 8.9 gap; top priority).** Suggested shape for the Schedule of Defaults: moderators are **confirmed by member vote** (simple majority) on appointment and at each annual renewal; any 10% of members may trigger a **recall vote**; the board may suspend immediately for cause pending a member vote. Keeps the legal floor's compliance duties intact while making stewards accountable downward, not upward.
- **P2 • Tenure-weight ratification with a sunset.** Put weighting to the promised cohort vote explicitly, and — whatever passes — attach a **re-ratification clause** (e.g., every 2 years) so a seniority aristocracy can never become furniture. Default-if-silent should be **flat**, matching the app's launch copy.
- **P3 • A restorative path in moderation.** A voluntary facilitated-conversation option for member-member conflict, alongside (never replacing) report/remove/appeal and never for legal-floor violations. Build post-launch; only then reflect it in the Terms (promise ≤ practice).
- **P4 • Hardship waiver by self-attestation.** One checkbox, no documentation, reviewed only in aggregate (count published in the transparency report).
- **P5 • Rules as visible objects (post-launch build).** A `/protected/rules` surface rendering the current Schedule of Defaults with each rule's amendment history and a “propose a change” affordance — modular politics made tangible.
- **P6 • Founder-protection visibility (post-launch build).** The transparency page states the protection period's terms and its countdown, in plain language, until it sunsets.

Explicitly not changed

Schema, RLS, auth, moderation code, governance numbers, and the Terms — the first four by standing constraint, the numbers because they are the cohort's to ratify, and the Terms because none of the above exists in the product yet.

Part IV --- Deliberate vs emergent: Rumelt × brown

The strategy corpus leans deliberately toward **Rumelt's Good Strategy** — diagnosis, guiding policy, coherent action. The live question is whether that contradicts **brown's emergent strategy** (Parts II-III). The answer, verified against the documents rather than asserted: no — and not because they “govern different layers,” but because **at Steppe they converge at the kernel**. (A side note for precision: Mintzberg's unrelated “emergent strategy” — plans as patterns learned in action — is a different doctrine; it supplies one useful mechanism below, nothing more.)

The Rumelt scorecard (the strategy corpus is genuinely “good strategy”)

Kernel element	Where Steppe has it	Evidence
Diagnosis	Extractive, engagement-optimized platform monoculture has hollowed local civic fabric; the honest hard problem is named — adoption (~2,300 paying members ≈ 6.5% of Redmond), not arithmetic	GTM Framing; BP v12; budget v1
Guiding policy (channels effort by excluding)	Trust transferred through real people is the only channel; dues-fund-core-always; structure-is-the-hook; the four non-negotiables rule whole tactic classes out permanently	GTM non-negotiables
Leverage from asymmetry	“A competitor can match a feature; it cannot match an ownership structure without restructuring” — transparency and ownership as the un-copyable advantage	GTM Framing ¶13
Proximate objectives	50-100 on the list <i>before</i> the public moment; a real vote in week one; funnel-before-push (“fix the pipe, then open the tap”)	GTM Phases 0-1
Coherent action	Vouching growth ↔ verification model ↔ transparency marketing ↔ captains structure — each reinforces the others; phases sequenced with hard dependencies	GTM Phases 0-5

Bad-strategy markers (fluff, goals-as-strategy, incoherent objectives, ducking the challenge): substantially absent — the remaining instances (mis-sized asks, the \$2/\$4 incoherence) were exactly what the funder-pack corrections removed.

The convergence: the guiding policy IS brown's doctrine in Rumelt's form

Look at what the kernel actually says. Rumelt demands a guiding policy that concentrates force on an asymmetry competitors can't copy. Steppe's guiding policy — *trust transferred through real people as the only channel; the way it grows is itself proof of what it is* — is **“move at the speed of trust” and “small is all” stated as strategy**. The captains fractal is simultaneously the Rumelt leverage (an un-copyable structural advantage) and the brown practice (the small-scale pattern that produces the large-scale outcome). “What you pay attention to grows” is why the health signals measure participation and trust rather than reach. The corpus doesn't *balance* Rumelt against brown; it uses Rumelt's discipline to concentrate effort on brown's mechanism. That coincidence is rare and worth protecting — most orgs' deliberate strategy fights their relational practice.

Where they divide labor: a fixed kernel with governed emergence

Rumelt treats strategy as a **hypothesis** to be tested and revised; brown makes adaptation a continuous, relational discipline rather than an annual replan (Mintzberg's pattern-promotion is the

mechanical version of the same loop). Steppe’s canon already resolves this the same way at every other layer — an invariant core with an adaptive periphery:

Layer	Fixed kernel (deliberate)	Governed emergence (adaptive)
Constitution	Non-votable legal floor; entrenched articles	Member-amendable Schedule of Defaults
Product	The invariants (no ads, no ranking, verify-then-forget)	Features, by member vote
Strategy	Diagnosis + guiding policy + non-negotiables	Tactics, channels, phase mechanics — hypotheses tested by the health signals

So the rule, now explicit (**ADOPTED A4**, codified in GTM v1.2): the strategy *kernel* changes only the way the constitution does — deliberately, in the open, with the members. The *action layer* is licensed to emerge, in brown’s sense: adaptation is continuous and relational (the A1 pulse and retro are where the felt experience of members enters the plan), with the health signals as the falsifiers Rumelt would demand — a tactic that fails its signals is wrong, and gets revised in a version bump, never patched silently. “Congruence is the test” polices the boundary — emergence never erodes the kernel, and the kernel never freezes the tactics.

One caution each way: pure Rumelt hardens into plan-worship (the map outliving the territory between rewrites); emergence without the kernel’s discipline drifts into the tyranny of structurelessness at the strategy layer — the same Freeman failure the captains model guards against. brown’s own guardrail (“intentional adaptation” — intentional first) and Rumelt’s (strategy as testable hypothesis) point at the same layered rule.

Changelog - v1.2 (2026-07-04) — Part IV re-centered on brown per clarification: the question was Rumelt × brown (not Mintzberg). Names the convergence — Steppe’s Rumelt guiding policy is brown’s doctrine in strategy form — and keeps the kernel/emergence layering rule (A4) with brown’s relational adaptation as the action-layer discipline. - v1.1 (2026-07-04) — Part IV added: Rumelt scorecard and the layering rule (A4) codified into GTM v1.2. - v1 (2026-07-04) — initial assessment; A1-A3 adopted; P1-P6 tabled for the cohort/board.